

The Queta Problem

Why Neemias Queta Is a Top-50 Player, and Why 2026–27 Proves It

GRAVITY Article No. 2 — a registered forecast

July 10, 2026 · revised July 11, 2026 (v2)

Not a research paper. This is an article, typeset in L^AT_EX.

Independent analysis, not peer reviewed. Findings may be revised. Errors are the author’s own.

Abstract

Our July 2026 ranking of all 644 NBA players (GRAVITY) placed Neemias Queta — undrafted-adjacent, twice-waived, the 39th pick of 2021 — at **#41**, ahead of household names and a full tier above his public reputation. This article treats that placement as a testable scientific claim rather than a quirk. We audit it three ways: we discount the model’s own shiniest input (his +8.6 on/off) with an empirical persistence model built from 184 player-pairs — read as uniform noise the ranking is unchanged, and under the harshest reading (only his number inflated) he slips just outside, to #53–55; we run a comparables engine over every statistically similar big-man season since 2020 (33 seasons from 22 players — the Gobert / Robert Williams / Jarrett Allen family) and find a top-50-caliber follow-up rate of 39% by season and 32% by player, before adjusting for Queta’s unusually secure situation; and we stress the one real threat, Boston’s signing of Mitchell Robinson — noting that the Celtics answered it themselves by extending Queta for \$56 million *afterward*. The article ends with a falsifiable forecast. GRAVITY prices the scenarios; the scenario weights are ours, anchored to observable rates in Section 7. Under those weights we assign roughly **75%** to Queta finishing 2026–27 as a top-50 player, and roughly **40%** to top-25. We will score this prediction in July 2027.

1 The claim, and why it sounds wrong

Every ranking system produces a few outputs its authors flinch at. In our July 2026 edition of “The 644,” the flinch was #41: Neemias Queta, ranked above Ivica Zubac, Rudy Gobert, Jaren Jackson Jr., and roughly four hundred players with larger reputations. The public model of Queta is a hustle big — a two-way-contract success story, fifth big on a title roster, fine. The statistical model of Queta is a starting center on a 56-win team who posted one of the most efficient high-volume finishing seasons in basketball, protected the rim, and led all Boston regulars in filtered on/off impact.

When a model and the crowd disagree this sharply, one of them is wrong, and articles that only ever defend their model are advertising. So this note does three things: it explains where #41 comes from, it attacks the number with the two strongest counterarguments we know how to formalize (noise and regression), and it converts what survives into a prediction with a scoring date. One relevant institution has already taken a side. On July 6, Boston signed Mitchell Robinson — a direct competitor for Queta’s minutes. Days later, they signed **Queta to a \$56 million extension anyway**.

2 How the rating works (the short version)

GRAVITY asks two questions: *how good is a player per minute on the floor* (box impact, efficiency relative to offensive burden, creation, defense, and garbage-time-filtered on/off from Cleaning the Glass), and *how much can you count on him* (games, role size, documented injuries). The two multiply. Everything is expressed in z-scores against the league-average minute: 0 average, +1 clearly good, +2 elite; the final scale sets 50 = average, 60+ = quality starter, 80+ = franchise engine (Jokić tops 2025–26 at 97.8). The full specification is in Appendix A; the seven-season historical variant (GRAVITY-RS) drops the on/off term, which is not uniformly available for early seasons.

3 Who is Neemias Queta, statistically

The first Portuguese-born NBA player; the 39th pick in 2021 out of Utah State; two two-way contracts; waived by Sacramento; a garbage-time cameo on Boston’s 2024 title team. Then Porziņģis was traded, Horford and Kornet left, and in 2025–26 coach Joe Mazzulla handed him the starting job: 76 games, 75 starts, 25.3 minutes per game, 10.2 points, 8.4 rebounds, 1.3 blocks — career highs across the board.

Table 1: Queta’s NBA career. IMPACT is per-minute value in z units (GRAVITY-RS, the no-on/off variant, league-standardized within each season); Pct is his percentile among 1,000+ minute players; Ctr-Pct among centers only. MP is total Basketball-Reference minutes; where CTG figures are cited elsewhere they are garbage-time-filtered and lower (768 in 2024–25, 1,924 in 2025–26).

Season	Team	G	MP	TS%	WS/48	BPM	IMPACT	Pct	Ctr-Pct
2021–22	SAC	15	120	.495	.051	−5.1	−0.46	24	9
2022–23	SAC	5	29	.607	.149	−3.4	−0.10	46	13
2023–24	BOS	28	333	.662	.259	+1.5	+0.50	77	61
2024–25	BOS	62	863	.674	.179	−0.7	+0.03	54	28
2025–26	BOS	76	1926	.674	.224	+2.8	+0.85	88	77
2026 playoffs	BOS	7	152	.767	.230	+1.2	—	—	—

Two mechanical details explain why the role took five years and why it should persist. First, the classic young-big constraint — fouling — has receded on schedule: his foul rate went from the 6th percentile in 2023–24 (a foul machine, unplayable in long stretches) to the 24th, to the 31st. Coaches could not keep him on the floor before; in 2025–26 they could. Second, his efficiency did not bend as the load grew.

In plain terms: Queta finally got starter minutes at 26 and his efficiency — which usually falls when a bench big’s role doubles — went up instead. He finishes like a top-10 center, crashes the offensive glass, blocks shots, and stopped fouling himself off the floor. That’s the whole trick: there is no empty-calorie stat propping him up.

4 Stress test I: shrink the shiny number

The skeptic’s first target should be that +8.6 on/off — Boston was 8.6 points per 100 possessions better with Queta on the floor (garbage time excluded). On/off numbers are notoriously unstable,




Figure 1: Five seasons: per-minute impact (top) and the role that finally arrived (bottom). The 2024–25 dip is real — his box value that year was ordinary — and the article does not hide it.

Table 2: The 2025–26 profile in CTG percentiles (among bigs). This is what #41 is made of: elite finishing at real volume, elite offensive rebounding, real rim protection, and the league’s 17th-best filtered on/off among 1,500+ minute players.

Component	Value	Percentile
Scoring efficiency (PSA)	134.0	84th
eFG%	65.4%	86th
Rim accuracy	75%	85th
Offensive rebound rate	12.3%	79th
Block rate	2.7%	79th
Defensive rebound rate	19.9%	74th
Filtered on/off (raw)	+8.6	#17 of 158 league-wide
Turnover rate	12.6%	55th
Foul rate	4.4%	31st (up from 6th in '24)
Free-throw %	70.1%	41st

so we measured the instability rather than argue about it. For all 184 players with 1,000+ CTG minutes in both 2024–25 and 2025–26, we correlated their on/off across the two seasons (Figure 3):

$$r = 0.317. \quad (1)$$

Call this what it is: *cross-season persistence*. Only about a third of the cross-player variation in filtered on/off carries into the following year. (Some of the rest is measurement noise; some is genuine change in teammates, role, and scheme — both argue for discounting any single season’s value.) Two technical notes a careful reader deserves. First, 0.317 is the Pearson correlation, but the season-to-season standard deviations are nearly equal (5.66 vs. 5.74), so the least-squares slope, $b = r s_{26}/s_{25} = 0.322$, is numerically the same discount. Second, there are two defensible shrink targets. The textbook regression prediction shrinks toward the qualifying pool’s mean (+1.2):

$$\widehat{\text{on/off}} = 1.18 + 0.317(8.6 - 1.18) = +3.5. \quad (2)$$




Figure 2: The anti-fluke chart. CTG scoring efficiency (PSA) against role size, Boston regular seasons (labels show total bbref minutes). The normal pattern is a downward slope — bench bigs feast on easy minutes and regress when the role doubles. Queta’s minutes per game roughly doubled and his season workload grew nearly sixfold (333 → 1,926 minutes), and the efficiency line went *up*: 130 → 132 → 134, all 79th–87th percentile. The open diamond is the seven-game 2026 playoff run (152 minutes, 97th percentile): a different competition environment on a small sample, plotted separately and no part of the regular-season trend.

The harsher convention shrinks toward zero:

$$\widehat{\text{on/off}}_{\text{harsh}} = r \times 8.6 = +\mathbf{2.7}. \quad (3)$$

We test the claim against both, and quote the harsher one in the summary.

This is the article attacking its own best exhibit, so state the result plainly — and the attack draws real blood. There are two coherent readings of “on/off is mostly noise,” and they cut very differently. *If every player’s on/off is equally unstable*, the discount changes nothing: GRAVITY uses on/off as a within-season z-score, and shrinking an entire column toward its mean leaves every z-score — and every rank — exactly where it was. #41 is untouched. The skeptic’s real claim must therefore be sharper: that *Queta’s number specifically* is inflated relative to his peers’. We tested that harshest reading by forcing his on/off input alone down to the discounted values and re-running the entire 644-player pipeline. At the regression value (+3.5) he lands at 57.2, **#53**; at the toward-zero value (+2.7), 56.9, **#55**. (Version 1 of this article reported $\approx$ #48 here, from a shortcut recomputation; the full re-run is the correct number, and it is worse for us.) So the honest summary: under the harshest single-player discount, Queta slips just *outside* the top 50 — the placement is carried mostly, but not entirely, by the box profile of Table 2 (on/off is 14% of the impact blend, Appendix A). The claim “top-50 player” narrows to the boundary. That is precisely why Section 7 registers the forward bet at 75% rather than declaring the present tense settled.

In plain terms: Only about a third of a season’s on/off carries into the next year — we measured it ($r = 0.317$, $n = 184$). If everyone’s number is equally noisy, #41 doesn’t move at all. If you believe Queta’s alone is inflated — the harshest reading — he lands #53–55, just outside the claim. The bet was never resting on the shiny number, but v2 owes you the truth that the harshest discount grazes it.




Figure 3: How much of an on/off number persists into the next year? The cloud says: about a third. The amber diamond is Queta — his 2024–25 on/off (a -0.8 in his backup season: 863 bbref minutes, 768 after CTG’s garbage-time filter, too few to qualify for the cloud) carried almost no information; his 2025–26 value is large but must be discounted by (2)–(3) before you believe it.

5 Stress test II: what happened to everyone like him

The second attack: *every* backup-turned-starter big has one shiny efficient season; most regress. Rather than argue by anecdote, we built the cohort. From six seasons of league data (2019–20 through 2024–25), we extracted every big-man season matching Queta’s 2025–26 statistical silhouette — usage ≤ 19 , WS/48 $\geq .190$, TS $\geq .630$, block rate $\geq 1.8\%$, 900–2,500 minutes, age ≤ 28 — and followed each into the next season. There are 33 such *seasons*, from 22 distinct players (a distinction the analysis below keeps honest), and the family is exactly who you’d hope: Gobert, Robert Williams, Jarrett Allen (four times), Zubac, Hartenstein, Gafford, Claxton, Kessler, Duren — and, three times, Mitchell Robinson.

Two robustness checks, because a critic should ask whether the criteria were tuned — consciously or not — to produce a flattering family. First, *players versus seasons*: repeat careers (Allen four times, Robinson three) give some players multiple votes. Counting each player once drops the top-50-caliber follow-up rate from 39% to 32%; clustering outcomes by player gives the same answer. The honest base rate is “roughly one in three,” and we carry both figures forward. Second, *the cutoffs*: Table 4 re-runs the engine varying one criterion at a time.

The base rates give the skeptic real ammunition — the median comparable gave back about half his impact z , and minutes *shrank* more often than grew — and they give the bull real support: 32–39% of the cohort (by player and by season) posted a top-50-caliber next season, the rate survives every cutoff variation we tested, and the family’s floor is remarkably high. Two situational facts then push Queta above the cohort median. Most comparables were backups fighting for a role (Kornet, Jackson-Davis, Clarke); Queta is the *incumbent starter* on a 56-win team. And no comparable had his team commit \$56M in July. Base rate roughly one in three, situation-adjusted: higher — and Section 7 must now say by how much, and why.

Table 3: Selected comparables (of 33): elite-efficiency, low-usage big seasons and what happened next. Full criteria in text; IMPACT in z units, Pct = next-season percentile among 1,000+ minute players.

Year	Player	Age	MP	IMPACT	Next MP	Next IMPACT	Next Pct
2020–21	Robert Williams	23	985	+1.66	1804	+1.48	95
2024–25	Jalen Duren	21	2034	+1.02	1976	+1.60	97
2019–20	Rudy Gobert	27	2333	+1.12	2187	+1.38	94
2023–24	Daniel Gafford	25	1815	+0.89	1226	+1.32	95
2024–25	Jarrett Allen	26	2296	+1.26	1519	+0.88	88
2023–24	Isaiah Hartenstein	25	1896	+0.87	1590	+0.86	84
2021–22	Isaiah Hartenstein	23	1216	+1.34	1626	−0.16	42
2022–23	Walker Kessler	21	1703	+0.89	1493	+0.24	67
2022–23	Nic Claxton	23	2271	+0.96	2116	+0.38	73
2021–22	Mitchell Robinson	23	1848	+0.75	1591	+0.63	82
2019–20	Mitchell Robinson	21	1412	+1.02	853	+0.34	69
2020–21	Enes Freedom	28	1758	+0.39	411	−0.23	38

All 33 seasons: median next-season percentile **75th**; ≥ 80 th: **45%**; ≥ 85 th (top-50-caliber): **39%**

22 unique players, counting each once: ≥ 85 th **32%**, median percentile 73rd (clustered player-means agree)

median impact regresses $+0.89 \rightarrow +0.47$; median minutes change **−146**; 55% kept $\geq 90\%$ of their minutes, 42% grew

Table 4: Cohort sensitivity: one cutoff varied at a time. The follow-up rate never falls below 38%, and the strictest efficiency cut *raises* it — the criteria were not tuned to flatter the conclusion.

Variation	Seasons	Unique players	≥ 85 th pct rate	Median next pct
As published	33	22	39%	75
TS $\geq .620$ (looser)	33	22	39%	75
TS $\geq .640$ (stricter)	29	18	41%	78
WS/48 $\geq .175$ (looser)	41	25	39%	74
WS/48 $\geq .205$ (stricter)	20	13	55%	86
Minutes $\geq 1,100$	29	20	38%	75
Usage ≤ 20	35	22	40%	75
Block rate $\geq 1.5\%$	33	22	39%	75

In plain terms: We found the 33 seasons (22 players) statistically most like 2025–26 Queta over the last six years. Most gave a little back the next season — but roughly one in three players (two in five season-weighted) delivered a top-50-level year, the rate survives every threshold we varied, and unlike most of them, Queta enters next season with the starting job, a new contract, and a 56-win team around him.

6 The Robinson question

The strongest objection isn’t statistical, it’s transactional: on July 6 Boston signed Mitchell Robinson — the champion Knicks’ rim protector — for three years and \$47.4M. If Robinson takes 24 minutes a night, the bear case writes itself.

Three replies, in declining order of hardness. *First*, the revealed preference: Boston extended Queta for \$56M days *after* the Robinson deal, and beat reporting projects Queta to open as the starter, with something like an even split behind him. Teams do sometimes pay two centers for depth, insurance, or matchups, so the contracts do not fix his role — what they establish is narrower but real: Boston values Queta at starter money while holding better information than any public




Figure 4: Regression to the mean, drawn honestly. Most dots sit below the dashed line: seasons like this usually give some value back. But the floor is high — three-quarters of the cohort remained above the league median, and two in five delivered a top-50-caliber (85th+ percentile) follow-up.

model, and plans to play him. *Second*, Robinson’s own availability record is the worst in the comparables family — 60+ games in only three of his last seven seasons, and he appears in Table 3 three times precisely because he keeps posting elite fragments (853, 768, 1,591 minutes follow-ups); planning a center rotation around Robinson’s health has broken repeatedly for a decade. *Third*, the tide: Jayson Tatum returns for a full season. His early post-Achilles evidence is encouraging — 16 regular-season games plus a seven-game playoff series at an 8.9 BPM — though samples that size cannot yet certify the recovery. It raises every Celtic’s context — and Queta’s playoff series alongside him (.767 TS) was his best basketball of the year. The honest version of the Robinson signing is a minutes tax on the bear case, not a refutation: it lowers Queta’s floor scenario from “starter by default” to “starter by performance,” which — given Section 5 — is a bet the data supports.

7 The prediction

Now the arithmetic — with the division of labor stated plainly, because it is the part of this article most deserving of scrutiny. The two-year GRAVITY machinery (Appendix A) converts any assumed 2026–27 performance level and workload into a score; the top-50 cutoff in the July 2026 distribution is **57.4**. *Those scores are model output*. The four scenarios and their weights are *not*: they are editorial judgments, and the subsection below anchors each one to an observable rate so the reader can price the judgment rather than take it.

Where the weights come from

The move from a one-in-three base rate to a 75% forecast is the least-supported leap a skeptic will find in this article, so here is each weight against the observable that disciplines it. **Bust, 15%:** losing the job outright requires Mitchell Robinson to be simultaneously healthy and clearly better — Robinson has reached 60 games in three of his last seven seasons ($\approx 43\%$), and the incumbent he must displace was just extended at starter money; the joint event deserves a weight well below

Table 5: Four futures, weighted. Impact is 2026–27 per-minute value in z units — his 2025–26 was +0.97 under current GRAVITY (which includes on/off) and +0.85 under GRAVITY-RS (the no-on/off variant used in Table 1 and the comparables engine); the blend automatically carries 2025–26 forward as the trailing season. Scores: model. Weights: authors.

Scenario	Weight	Impact	MP	GP	Score	Top-50?	Reading
Bust: loses the job	15%	+0.30	1000	50	53.8	no	Robinson takes over
Bear: even timeshare	25%	+0.75	1450	62	58.1	<i>barely</i>	regression + split
Base: holds the job	40%	+1.00	2000	74	62.3	yes (#23–27)	2025–26 continues
Bull: consolidation	20%	+1.25	2300	76	66.9	yes (#10–13)	comps’ top tercile

$P(\text{top-50}) \approx 0.25 \times 0.6 + 0.40 + 0.20 = \mathbf{0.75}$ (bear counted 60%, being borderline — a judgment, like the weights)
 $P(\text{top-25}) \approx 0.40 \times 0.5 + 0.20 = \mathbf{0.40}$




Figure 5: The bet, drawn. The dashed amber line is the top-50 cutoff in the July 2026 distribution (57.4). Three of four scenarios clear it; under our weights the assigned probability is about 75%.

either margin. **Bear, 25%:** the cohort says roles shrink more often than they grow (median −146 minutes; only 42% grew), so a timeshare is the modal downside and gets the second-largest weight. **Base, 40%:** 55% of comparables kept at least 90% of their minutes, and Queta holds two advantages most of them lacked — incumbency and the contract — so “2025–26 repeats” is the single most likely outcome, weighted slightly below the naked retention rate out of respect for Robinson. **Bull, 20%:** one third of the cohort landed in the top tercile of follow-ups; we haircut that for the crowded frontcourt. These anchors discipline the weights; they do not compute them. A reader who rejects every situational adjustment should simply hold the cohort rate — roughly one in three — and the distance between that number and 75% is exactly the size of the bet this article places on incumbency, the extension, Robinson’s availability record, and Tatum’s return.

Note also what makes the claim robust: it does not require a leap. The base case is simply *2025–26 happening again* — same job, same efficiency, ordinary regression — and it lands at #23–27, comfortably inside the claim. The prediction fails only if Robinson takes the job outright or the season-scale regression lands in the worst tercile of the comparables. **Falsification criteria, on the record:** in the July 2027 edition of “The 644,” computed by this same specification, Queta at #50 or better scores as a win; below #70 as a clean miss. The Ledger’s registered middle band (#51–70, originally labeled a “push”) stands as registered, but v2 tightens the grading against ourselves: *the claim is “top-50,” so #51 is a failure*. A #51–70 finish will be graded a **near miss** — a failed prediction owing a diagnosis, not partial credit — and the binary verdict, top-50 or not, is what the July 2027 audit will headline.

In plain terms: The bet: we assign about 75% to Queta being a top-50 player in next July’s ranking, and 2-in-5 to top-25 — the model prices the scenarios, we chose the weights, and the anchors above show our work. It’s not a moonshot claim; it just requires this season to mostly repeat. We wrote down in advance exactly what counts as being wrong, and #51 counts.

8 Conclusion

The Queta problem narrows under measurement to something bettable. A #41 ranking that looks like a model glitch turns out to rest on the most trustworthy statistics in the sport — hundreds of finished possessions at 85th-percentile efficiency, elite offensive rebounding, real rim protection — and the two best attacks we could mount on it leave marks without landing: discounting his on/off by its measured cross-season persistence ($r = 0.317$) leaves #41 untouched if the discount applies to everyone, and puts him at #53–55 under the harshest reading that his number alone is inflated; the 33-season, 22-player comparables cohort puts the base rate of a top-50-caliber follow-up at roughly one in three, robust to every cutoff we varied, before counting his incumbency, his new contract, or Tatum’s return. The model prices the futures; the weights that turn those prices into a 75% forecast are ours, anchored and argued in Section 7. Boston’s own front office, holding better information than any public model, paid \$56 million for a related conclusion. The market will grade us in twelve months.

Revision note (v2, July 11, 2026)

Following a detailed external review, this version separates model output from editorial judgment and corrects one number. (1) The scenario probabilities were presented as GRAVITY output; they are the authors’ weights applied to GRAVITY scores, and a new subsection anchors each weight to an observable rate. (2) “On/off is two-thirds noise” overstated what a cross-season correlation shows; the language now says *persistence*, reports the regression slope alongside r (they are numerically equal here), and gives both shrink targets. (3) The shrunk-on/off counterfactual was corrected: a full re-run of the 644-player pipeline gives #53–55, not the $\approx$ #48 a shortcut recomputation produced in v1 — the harshest single-player discount moves Queta just outside the top 50, and the article now says so; the uniform-discount case (ranks unchanged) was added for the same honesty. (4) The comparables cohort now discloses 33 seasons from 22 unique players, reports per-player rates (32% vs. 39%), and adds a cutoff-sensitivity table. (5) The grading of the registered forecast was tightened against ourselves: #51–70 will be graded a near miss, not a push. (6) Figure 2 now labels minutes and separates the playoff sample from the regular-season trend; minutes sources (bbref vs. CTG filtered) are labeled throughout; the Robinson-contract and Tatum-recovery inferences were softened to what the evidence supports. The registered forecast itself — 75% top-50, 40% top-25, scored July 2027 — is unchanged: that is the bet, and moving it after registration is the one revision this publication does not allow.

Relabel note (October 6, 2026)

This document was first published as a research note. It is relabeled as an article: our own analysis, typeset in L^AT_EX, not a research paper and not peer reviewed. No numbers, findings, or forecasts changed.

A The GRAVITY specification (compact)

Identical to Article No. 1 and the July 10, 2026 “The 644” ranking (v3). For season y , minutes-weighted moments define clipped z-scores $z_k(i) = \text{clip}((x_k(i) - \mu_k)/\sigma_k, \pm 4)$. Components: efficiency-under-load $(\text{TS} - \mu_{\text{TS}}) \cdot 100 \cdot (\text{clip}(\text{USG}, 8, 36)/20)^{0.7}$; creation $\text{AST}\% - 0.55 \text{TOV}\%$; load $\text{USG} + 0.4 \text{AST}\%$; stocks $2 \text{STL}\% + 1.2 \text{BLK}\%$; rim pressure; volume (pts/100). Skill blend $S = 0.27 z_{\text{vol}} + 0.19 z_{\text{eff}} + 0.22 z_{\text{crea}} + 0.09 z_{\text{load}} + 0.10 z_{\text{stk}} + 0.05 z_{\text{reb}} + 0.08 z_{\text{rim}}$. Season impact

$$I_y = [0.44 z_{\text{BPM}} + 0.12 z_{\text{WS}/48} + 0.14 z(\Delta_{\text{on/off}} \min(1, m/2400)) + 0.30 S] \sqrt{\min(1, m/800)}, \quad (4)$$

blended across two seasons with minute-proportional weights (0.75 on the recent year), shrunk toward replacement $\rho = -1.9$ with credibility $c = \sqrt{\min(M, 2200)/2200}$, adjusted for playoff evidence and age, then discounted by availability $A = 0.78 + 0.22 \min(1, (0.7 g_y + 0.3 g_{y-1})/55)$, role size $\Phi = 0.62 + 0.38 \min(1, \text{mpg}^*/32)$, and documented injuries $J \in [0.78, 1]$. Final scale $G = 50 + 15[\rho + (\tilde{R} + \pi + \alpha - \rho)A\Phi J]$. The historical variant GRAVITY-RS (used for Table 1 and the comparables engine) deletes the on/off term and renormalizes to $[0.533 z_{\text{BPM}} + 0.133 z_{\text{WS}/48} + 0.333 S'] \sqrt{\min(1, m/800)}$.

Data and forecast registration

Basketball-Reference league tables 2019–20 through 2025–26 ($\approx 4,400$ player-seasons); Cleaning the Glass subscriber database (filtered on/off, shot zones, foul and rebounding rates; player page 4902), accessed July 10, 2026. The on/off cross-season persistence estimate uses all 184 players with 1,000+ CTG minutes in both 2024–25 and 2025–26. No external rankings or projections were consulted. Prediction registered July 10, 2026; to be scored against the July 2027 “The 644.”

References

- [1] Basketball-Reference.com, league tables, accessed July 10, 2026.
- [2] Cleaning the Glass, player and leaders tables (subscription). <https://cleaningtheglass.com/stats/player/4902>
- [3] “Why Neemias Queta is the Celtics’ starting center for 2025–26,” Yahoo Sports / NBC Sports Boston, October 2025.
- [4] “Celtics sign Neemias Queta to \$56 million extension after Mitchell Robinson addition,” Yahoo Sports, July 2026.
- [5] “Mitchell Robinson joining Celtics on 3-year, \$47.4M deal,” NBA.com / Hoops Rumors, July 6, 2026.
- [6] GRAVITY Article No. 1, “The Rise and Fall of Ja Morant,” July 10, 2026.